AI Usage and Governance Policy

Document ID: AI-008 | Owner: AI Governance | Effective date: 2026-04-01 | Version 1.0

This policy governs the development, deployment, and use of artificial intelligence solutions at Nile Horizon Bank, including generative AI and large language models (LLMs). It is a fictional policy for training purposes.

1. AI Use Case Registration and Risk Tiers

Every AI use case must be registered in the model inventory before development starts.

Use cases are assigned a risk tier. A use case is High risk if it influences customer eligibility, pricing, or credit decisions. Use cases that only support internal staff productivity are normally Low risk.

Risk tier	Approver
Low	AI Engineering Manager
Medium	AI Governance Committee
High	AI Governance Committee and Risk Management

2. Data Rules for Generative AI

Restricted data must not be sent to externally hosted LLMs under any circumstances.

Confidential data may only be sent to approved providers that have a contract guaranteeing zero data retention. Personal data must be masked in prompts and in logs.

3. Human Oversight

High risk AI solutions require a human in the loop who reviews and can override the output before any customer is affected.

Knowledge assistants must show the source documents used for each answer, and answers must be clearly labelled as AI-generated.

4. Testing Before Release

Every AI service must pass unit tests and integration tests before release. Customer-impacting models must also be tested for bias and fairness. Services that accept free-text input must be red-teamed for prompt injection before release.

5. Monitoring and Drift

Performance and data drift must be reviewed every month. If the Population Stability Index (PSI) of a monitored input exceeds 0.2, the owner must start an investigation within 5 business days.

Significant issues must be escalated to the AI Engineering Manager.

6. Versioning and Change Control

Models, prompts, and indexes used in production must be versioned. Any change in production requires an approved change request and a documented rollback plan.

7. Incidents and Documentation

Any AI incident, such as a harmful or incorrect output reaching a customer, must be reported to the AI Engineering Manager within 24 hours.

Each production AI service must have a model card, an API specification, and an operational runbook.